


R data types: Lists and Data Frames

Alberto Garfagnini

Università di Padova

R lecture 3



R internals: variables and objects creation

- We create a vector with three values and assign it to a reference variable, `x`

```
x <- c(1,2,3)
```

- we now copy `x` to another variable `y`:

```
y <- x
```

- and modify one element of `y`

```
y[3] <- 4
```


- did we modify also `x`?

No, they refer to two different objects:

```
str(x)
%> num [1:3] 1 2 3
str(y)
%> num [1:3] 1 2 4
```

- the behavior is called **copy-on-modify**
- all R objects are **immutable**

```
lobstr::obj_addr(x)
"0x55d03cd66fb8"
```



```
lobstr::obj_addr(y)
"0x55d03dbac8c8"
```



The `lobstr` package allows to visualize R data structures: it shows memory location and size of objects.

URL: <https://github.com/r-lib/lobstr>

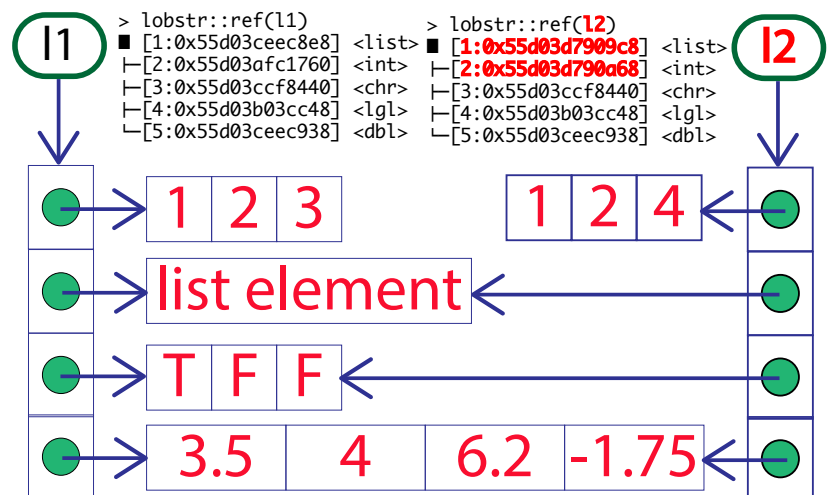
R Lists

- **Lists** are an evolution of atomic vectors: **each element can be of any type**
- from the technical point of view: **each element of a list** is of the same type: it is **a reference to another R object**
- building a list:

```
l1 <- list( 1:3,
            "list_element",
            c(TRUE, FALSE, FALSE),
            c(3.5, 4, 6.2, -1.75)
)
typeof(l1)
%> [1] "list"
```

- we copy to a new list and modify one element

```
l2 <- l1
l2[[1]] <- c(1L, 2L, 4L)
```



R matrices

- a matrix is a 2-dimensional object
- the first way of creating a matrix is by calling the `matrix()` object constructor

```
X <- matrix(c(1,0,0,0,1,0,0,0,1), nrow=3) ; X
%>      [,1] [,2] [,3]
%> [1,]    1    0    0
%> [2,]    0    1    0
%> [3,]    0    0    1

class(X)
%> [1] "matrix"
attributes(X)
%> $dim
%> [1] 3 3
str(X)
%> num [1:3, 1:3] 1 0 0 0 1 0 0 0 1
```

- another way is to transform a vector in a matrix: data can be arranged by rows (`byrow=T`) or columns (`byrow=F`)

```
vct <- c(1,2,3,4,4,3,2,1)
V <- matrix(vct, byrow=T, nrow=2)
V
%>      [,1] [,2] [,3] [,4]
%> [1,]    1    2    3    4
%> [2,]    4    3    2    1
```

```
V <- matrix(vct, byrow=F, nrow=2)
V
%>      [,1] [,2] [,3] [,4]
%> [1,]    1    3    4    2
%> [2,]    2    4    3    1
```

- another possibility is to convert the vector to a matrix by specifying the new dimensions (rows and columns), using the `dim` function

```
vct <- c(1,2,3,4,4,3,2,1)
vct
%> [1] 1 2 3 4 4 3 2 1
```

```
dim(vct) <- c(4,2)
is.matrix(vct)
%> [1] TRUE
```

```
vct
%>      [,1] [,2]
%> [1,]    1    4
%> [2,]    2    3
%> [3,]    3    2
%> [4,]    4    1
```

- we can then transform the matrix:

```
tvct <- t(vct) # transpose the matrix
tvct
%>      [,1] [,2] [,3] [,4]
%> [1,]    1    2    3    4
%> [2,]    4    3    2    1
```

Accessing or operating on matrix rows or columns

- Let's create a matrix with $n = 20$ entries sampled from a Poisson distribution with $\lambda = 1.5$

```
X <- matrix(rpois(n=20,lambda=1.5), nrow=4)
X
%>      [,1] [,2] [,3] [,4] [,5]
%> [1,]    1    1    1    2    4
%> [2,]    1    1    3    3    2
%> [3,]    1    3    5    0    1
%> [4,]    2    1    1    2    2

X[3,3] # return element in row 3 and column 3
%> [1] 5
X[4,] # return row 4
%> [1] 2 1 1 2 2
X[,5] # return column 5
%> [1] 4 2 1 2
```

- there are special functions for calculating summary statistics on a matrix:

```
rowSums(X) # use colSums(X) for columns
%> [1] 9 10 10 8
rowMeans(X) # use colMeans(X) for columns
%> [1] 1.8 2.0 2.0 1.6
```

Adding rows and columns to a matrix

- given a matrix, we would like to add a row, at the bottom, showing the column means, and a column at the right showing the row variances:

```
vct <- matrix(c(1,0,2,5,1,1,3,1,3,1,0,2,1,0,2,1), byrow=T, nrow=4)
vct
%>      [,1] [,2] [,3] [,4]
%> [1,]    1    0    2    5
%> [2,]    1    1    3    1
%> [3,]    3    1    0    2
%> [4,]    1    0    2    1

vct <- rbind(vct, apply(vct, 2, mean))
vct <- cbind(vct, apply(vct, 1, var))

colnames(vct) <- c(1:4, "variance")
rownames(vct) <- c(1:4, "mean")

vct
%>      1    2    3    4  variance
%> 1    1.0 0.0 2.00 5.00 4.6666667
%> 2    1.0 1.0 3.00 1.00 1.0000000
%> 3    3.0 1.0 0.00 2.00 1.6666667
%> 4    1.0 0.0 2.00 1.00 0.6666667
%> mean 1.5 0.5 1.75 2.25 0.5416667
```

apply(), sapply() and lapply()

- `apply()` is used to apply functions to rows or columns of matrices or dataframes across one of the margins of a matrix
- `margin=1` refers to the rows and `margin=2` to the columns

```
(Y <- matrix(rbinom(20, 9, 0.45), nrow=4))
%>      [,1] [,2] [,3] [,4] [,5]
%> [1,]    6    5    4    2    5
%> [2,]    6    3    3    5    4
%> [3,]    5    2    3    4    4
%> [4,]    3    4    4    3    6
apply(Y, MARGIN=2, FUN=sum) # apply sum() to all columns
%> [1] 20 14 14 14 19
```

- we can `apply()` functions to the individual elements of a matrix. In this case, the `margin` parameter, determines only the shape of the resulting matrix

```
apply(Y, 1, sqrt)
%>      [,1] [,2] [,3] [,4]
%> [1,] 2.449490 2.449490 2.236068 1.732051
%> [2,] 2.236068 1.732051 1.414214 2.000000
%> [3,] 2.000000 1.732051 1.732051 2.000000
%> [4,] 1.414214 2.236068 2.000000 1.732051
%> [5,] 2.236068 2.000000 2.000000 2.449490
apply(Y, 2, sqrt)
%>      [,1] [,2] [,3] [,4] [,5]
%> [1,] 2.449490 2.236068 2.000000 1.414214 2.236068
%> [2,] 2.449490 1.732051 1.732051 2.236068 2.000000
%> [3,] 2.236068 1.414214 1.732051 2.000000 2.000000
%> [4,] 1.732051 2.000000 2.000000 1.732051 2.449490
```

- it is also possible to apply an anonymous, user defined, function

```
apply(Y, 1, function(x) x^2+x) # compute x^2 + x for each element
%>      [,1] [,2] [,3] [,4]
%> [1,]    42    42    30    12
%> [2,]    30    12     6    20
%> [3,]    20    12    12    20
%> [4,]     6    30    20    12
%> [5,]    30    20    20    42
```

- in case you need to apply a function to a vector, rather than to the margin of a matrix, use `sapply()`

```
sapply(12:14, seq) # generate a list of seq, from 1:12 to 1:14
%> [[1]]
%> [1]  1  2  3  4  5  6  7  8  9 10 11 12
%>
%> [[2]]
%> [1]  1  2  3  4  5  6  7  8  9 10 11 12 13
%>
%> [[3]]
%> [1]  1  2  3  4  5  6  7  8  9 10 11 12 13 14
```

Random numbers

- random numbers from a uniform distribution $\mathcal{U}(0, 1)$ are generated using `runif()`
- the random generation seed is set via `set.seed()`

```
set.seed(2019)
runif(3)
%> [1] 0.7699015 0.7128397 0.3033602
```

- resetting the seed with the same value will generate the same sequence of random numbers
- it is also possible to save the current seed and reuse it to obtain the same random numbers sub-sequence

```
current.seed <- .Random.seed # save the current seed
runif(3)
%> [1] 0.61823636 0.05048374 0.04321880
runif(3)
%> [1] 0.820176206 0.009614496 0.102491504

current.seed -> .Random.seed # reset the previous sequence seed
runif(5)
%> [1] 0.618236361 0.050483740 0.043218804 0.820176206 0.009614496
```

sampling from a vector

- generating random numbers from probability distributions will be discussed in the next lessons
- now we want to randomize (shuffling or sampling from) the elements of a vector
- There are two ways of sampling:
 - 1) **sampling without replacement** : all the vector values will appear in output, but in a randomized sequence
 - 2) **sampling with replacement** : some vector values may be re-selected and appear more than once in the output
- using `sample()`, sampling without replacement is the default operation

```
y <- c(8,3,5,7,6,6,8,9,2,3,9,4,10,4,11)
sample(y) # reshuffling all vector values
%> [1] 3 9 2 5 4 8 6 8 6 4 10 3 7 11 9
sample(y, 5) # pick up only 5 values from the original vector
%> [1] 3 8 9 4 8
sample(y, 5) # just redo it, and a different sequence may appear
%> [1] 8 3 8 4 3
```

- The option `replace=T` allows for sampling with replacement

```
sample(y, replace=T)
%> [1] 8 3 6 8 8 4 3 7 10 9 10 9 4 4 7
```

`sample()`'s surprise example

```
x <- 1:10
x
%> [1] 1 2 3 4 5 6 7 8 9 10

sample(x[x>8])
%> [1] 10 9

sample(x[x>9])
%> [1] 1 10 8 7 6 5 4 2 9 3

sample(x[x > 10])
%> integer(0)
```

- the first argument of `sample()` can be a vector of more than one element or an integer
- the `re-sample()` function is safer

```
sample(15)
%> [1] 2 4 3 11 7 14 6 5 1 13 15 10 12 9 8

str(x[x>9])
%> int 10

resample(x[x>8])
%> [1] 10 9
resample(x[x>9])
%> [1] 10
```

`sample(x, size, replace = FALSE, prob = NULL)`
If 'x' has length 1, sampling takes place from '1:x'

- R's subsetting operators are fast and powerful, and allow to perform complex operations in a way that few other languages can match
 - there are 6 ways to subset atomic vectors
 - there are 3 subsetting operators: `[[`, `[` and `$`
 - subsetting can be combined with assignment

Subsetting atomic vectors - 1

```
x <- c(2.1, 4, 6.7, 1.75)
```

- **positive integers** return elements at a specified position

```
x[c(1,3)]  
%> [1] 2.1 6.7
```

```
% Duplicate indices will duplicate values  
x[c(1,1,3,3)]  
%> [1] 2.1 2.1 6.7 6.7
```

```
% Real numbers are truncated to integers  
x[sort(x)]  
%> [1] 2.10 4.00 1.75 NA
```

- **negative integers** exclude elements

```
x[-c(1,3)]  
%> [1] 4.00 1.75
```

```
% NB negative and positive ints cannot be mixed  
x[c(-1,3)]  
%> Error in x[c(-1, 3)]: only 0's may be mixed with negative subscripts
```

Subsetting atomic vectors - 2

```
x <- c(2.1, 4, 6.7, 1.75)
```

- **logical vectors** select elements where the logical value is **TRUE**

```
x[c(T, T, F, T)]
%> [1] 2.10 4.00 1.75
```

```
x[x>2]
%> [1] 2.1 4.0 6.7
```

- if in `x[sel]`, `length(sel) != length(x)` the **recycling rules** are used: the shorter vector is recycled to the length of the longer

```
> x[c(TRUE, FALSE)]
[1] 2.1 6.7
```

```
%# is equivalent to:
> x[c(TRUE, FALSE, TRUE, FALSE)]
[1] 2.1 6.7
```

- **nothing** returns the original vector

```
x[]
%> [1] 2.10 4.00 6.70 1.75
```

Subsetting atomic vectors - 3

```
x <- c(2.1, 4, 6.7, 1.75)
```

- **zero** returns a zero-length vector (it can be helpful to generate test data)

```
x[0]
numeric(0)
```

- **named vectors** can be accessed with **character vectors**

```
y <- setNames(x, LETTERS[1:length(x)])
y
%>      A      B      C      D
%> 2.10 4.00 6.70 1.75
y["A"]
%>      A
%> 2.1
```

```
y[c('A', 'A', 'D')]
%>      A      A      D
%> 2.10 2.10 1.75
```

- **WARNING:** subsetting with factors will use the underlying integer vector, not the character levels. → **Avoid subsetting with factors**

```
y[factor("B")]
%>      A
%> 2.1
```


Subsetting matrices

- subsetting a matrix or a list works in a similar way as subsetting atomic vectors

```
S <- matrix(1:9, nrow = 3)
%> [1,] 1 4 7
%> [2,] 2 5 8
%> [3,] 3 6 9
```

- using `[]` always returns a list
- `[[]` and `$` allows to pull out elements from the list
- the common rule to subset a matrix (2D) and an array (nD , $n > 2$) is to supply a 1D vector for each dimension, separated by a comma
- blank subsetting allows to keep all data for the corresponding dimension

```
%# Get rows 1 and 3 and all columns
S[c(1,3), ]
%>      [,1] [,2] [,3]
%> [1,]    1    4    7
%> [2,]    3    6    9

colnames(S) <- c("S1", "S2", "S3")
S[c(T, F, T), c("S1", "S3")]
%>      S1 S3
%> [1,]  1  7
%> [2,]  3  9
```

Subsetting matrices - 2

- matrices and arrays are just vectors with special attributes, therefore they can be subset with a single vector, as if they were a 1D vector

```
v <- outer(1:5, 1:5, FUN="paste", sep=",")
v
%>      [,1] [,2] [,3] [,4] [,5]
%> [1,] "1,1" "1,2" "1,3" "1,4" "1,5"
%> [2,] "2,1" "2,2" "2,3" "2,4" "2,5"
%> [3,] "3,1" "3,2" "3,3" "3,4" "3,5"
%> [4,] "4,1" "4,2" "4,3" "4,4" "4,5"
%> [5,] "5,1" "5,2" "5,3" "5,4" "5,5"

v[seq(3, 23, 5)]
%> [1] "3,1" "3,2" "3,3" "3,4" "3,5"
```

- to preserve the original matrix dimension, use `drop = FALSE`

```
(S <- matrix(1:6, nrow = 2))
%>      [,1] [,2] [,3]
%> [1,]    1    3    5
%> [2,]    2    4    6

S[1, ]
%> [1] 1 3 5

S[1, , drop = FALSE]
%>      [,1] [,2] [,3]
%> [1,]    1    3    5
```

Selecting a single element


- there are two other subsetting operators:
 - `[]` is used to extract single items
 - `$` is used as a shorthand: `x$y` stands for `x[["y"]]`
- `[]` is most important while working with lists: subsetting a list with single `[]` always returns a smaller list

If list `xl` is a train carrying objects, then `xl[[5]]` is the object in car 5; `xl[4:6]` is a train of cars 4-6


<https://twitter.com/RLangTip/status/268375867468681216>

- with this metaphor let's build a list


```
xl <- list(1:3, "one", c(T,F,F))
```




Selecting a single element






- two options are available when extracting a single element:
 - create a smaller train, with fewer cars (using `[]`)



- or extract the content of a particular car (with `[[]]`)



- extracting multiple (or zero) elements, we have to build a smaller train



Loops

- Let's create a function, using loops, to evaluate the factorial:
- $n! = n \cdot (n - 1) \cdot (n - 2) \dots 2 \cdot 1$

```
fac1 <- function(x) {  
  f <- 1  
  if (x<2) return(1)  
  for (i in 2:x) {  
    f <- f*i  
  }  
  return(f)  
}  
sapply(1:5, fac1)  
%> [1] 1 2 6 24 120
```

```
fac2 <- function(x) {  
  f <- 1; t <- x  
  while (t>1) {  
    f <- f*t  
    t <- t-1  
  }  
  return(f)  
}  
sapply(1:5, fac2)  
%> [1] 1 2 6 24 120
```

```
fac3 <- function(x) {  
  f <- 1; t <- x  
  repeat {  
    if (t<2) break  
    f <- f*t  
    t <- t-1  
  }  
  return(f)  
}  
sapply(1:5, fac3)  
%> [1] 1 2 6 24 120
```

- But it is almost always better to use a built-in function that operates on the entire vector, removing the need of loops or repeats

```
> cumprod(1:5) # it does not work for 0  
[1] 1 2 6 24 120  
> fac4 <- function(x) max(cumprod(1:x))  
> sapply(1:5, fac4)  
[1] 1 2 6 24 120  
  
> # R implements a factorial() function, introduced not long ago  
> sapply(1:5, factorial)  
[1] 1 2 6 24 120
```

Loop avoidance: use vectorized operations

- it's a good R programming practice to avoid loops wherever possible
- in many cases, using vector functions, makes it particularly straightforward

```
> y <- c(-3,4,-2,-1,8,7,9)  
> y  
[1] -3 4 -2 -1 8 7 9  
  
> for (i in 1:length(y)) {if (y[i] < 0) y[i] <- 0}  
> y  
[1] 0 4 0 0 8 7 9
```

- in the example below, a loop can be replaced by logical subscripts

```
> y <- c(-3,4,-2,-1,8,7,9)  
> y[y<0] <- 0  
  
> y  
[1] 0 4 0 0 8 7 9
```

the ifelse() vectorized function

- ifelse() allow to work on an entire vector without using loops

```
• > y <- log(rpois(20,1.5))
> y
[1]      -Inf  1.0986123  0.0000000  0.0000000  0.0000000  0.6931472
[7]  0.6931472      -Inf  1.3862944  0.6931472  1.3862944      -Inf
[13]      -Inf  0.0000000  0.0000000  1.0986123  0.0000000  0.0000000
[19]      -Inf  1.0986123

> mean(y)
[1] -Inf

> (y <- ifelse(y<0, NA, y))
[1]      NA  1.0986123  0.0000000  0.0000000  0.0000000  0.6931472
[7]  0.6931472      NA  1.3862944  0.6931472  1.3862944      NA
[13]      NA  0.0000000  0.0000000  1.0986123  0.0000000  0.0000000
[19]      NA  1.0986123

> mean(y, na.rm=TRUE)
[1] 0.5431911
```

Loops are slow, compared to vectorized operations

- let's generate $5 \cdot 10^7$ events according to an uniform distribution, $\mathcal{U}(0,1)$
- we want to search for the maximum value in the vector using the vectorized function `max()` e by using conventional loops

```
x <- runif(50000000)

system.time(max(x))
%>   user  system elapsed
%>  0.106   0.000   0.106

pc <- proc.time()
cmax <- x[1]
for (i in 2:length(x)) { if(x[i]>cmax) cmax <- x[i] }

proc.time()-pc
%>   user  system elapsed
%>  2.061   0.071   2.133
```

- `system.time()` and `proc.time()` produce a vector of three numbers, showing the user, system and total elapsed time in seconds


Good/Bad practice in building vectors

- we want to build a vector containing 10^n elements in the sequence $1:10^n$
- three ways are analyzed

```
test1 <- function(n){  
  y <- 1:n  
}  
  
test2 <- function(n){  
  y <- numeric(n)  
  for (i in 1:n)  
    y[i] <- i  
}  
  
test3 <- function(n){  
  y <- NULL  
  for (i in 1:n)  
    y <- c(y,i)  
}
```

```
> system.time(test1(10000000))  
   user  system elapsed  
0.006   0.000   0.006  
> system.time(test2(10000000))  
   user  system elapsed  
0.622   0.011   0.633  
> system.time(test3(10000000))  
   user  system elapsed  
2.755   0.003   2.758
```

- the first method (test1) is the best
- the loop using a pre-determined vector length is reasonably fast
- the last method (test3) is the slowest
- Moral: **never grow vectors by repeated concatenation**



R Lists example

- let's create a more complex list with heterogeneous object types

```
apples <- c(4, 4.5, 5, 3.9)  
oranges <- c(TRUE, FALSE, TRUE)  
chalk <- c("limestone", "marl", "ooline", "CaCO3")  
pears <- c(3.2-4.5i, 12.8+2.2i)  
  
items <- list(apples, oranges, chalk, pears)  
items  
%> [[1]]  
%> [1] 4.0 4.5 5.0 3.9  
%>  
%> [[2]]  
%> [1] TRUE FALSE TRUE  
%>  
%> [[3]]  
%> [1] "limestone" "marl" "ooline" "CaCO3"  
%>  
%> [[4]]  
%> [1] 3.2-4.5i 12.8+2.2i
```

R List example: element access

- vectors, matrices and arrays subscripts have one set of square brackets [6], [3,4] or [2,3,2,1]
- lists subscripts have double square brackets [[2]] or [[i,j]]

```
items[[3]]
%> [1] "limestone" "marl"      "ooline"      "CaCO3"
```

```
items[[3]][1]
%> [1] "limestone"
```

- if the list elements have names, it is possible to use the operator \$ for list indexing

```
names(items) <- c("apples", "oranges", "chalk", "pears")
```

```
str(items)
%> List of 4
%> $ apples : num [1:4] 4 4.5 5 3.9
%> $ oranges: logi [1:3] TRUE FALSE TRUE
%> $ chalk  : chr [1:4] "limestone" "marl" "ooline" "CaCO3"
%> $ pears  : cplx [1:2] 3.2-4.5i 12.8+2.2i
```

```
items$pears
%> [1] 3.2-4.5i 12.8+2.2i
```

R list example: Applying functions

- the length of the list is the number of items on the list. To get the length of the individual vectors we use the lapply() function

```
length(items)
%> [1] 4
```

```
lapply(items, length)
%> $apples
%> [1] 4
```

```
%> $oranges
%> [1] 3
```

```
%> $chalk
%> [1] 4
```

```
%> $pears
%> [1] 2
```

```
class(items)
%> [1] "list"
```

```
lapply(items, class)
%> $apples
%> [1] "numeric"
```

```
%> $oranges
%> [1] "logical"
```

```
%> $chalk
%> [1] "character"
```

```
%> $pears
%> [1] "complex"
```

R Lists : 3

- applying numeric functions to the list, will only work for objects of class `numeric` or `complex`

```
mean(items)
%> [1] NA
%> Warning message:
%> In mean.default(items) :
%> argument is not numeric or logical: returning NA

lapply(items, mean)
%> $apples
%> [1] 4.35

%> $oranges
%> [1] 0.6666667

%> $chalk
%> [1] NA

%> $pears
%> [1] 8-1.15i
%> Warning message:
%> In mean.default(X[[i]], ...) :
%> argument is not numeric or logical: returning NA
```

- a warning message points out that the third vector cannot be coerced to a number (it is not numeric, complex or logical), therefore NA appears in the output

R Lists : 4

- The `summary()` function works for lists, but the most useful overview of a list content is given by `str()`, the structure function:

```
summary(items)
%>      Length Class  Mode
%> apples    4      -none- numeric
%> oranges    3      -none- logical
%> chalk      4      -none- character
%> <NA>       2      -none- complex

str(items)
%> List of 4
%> $ apples : num [1:4] 4 4.5 5 3.9
%> $ oranges: logi [1:3] TRUE FALSE TRUE
%> $ chalk  : chr [1:4] "limestone" "marl" "ooline" "CaCO3"
%> $ NA     : cplx [1:2] 3.2-4.5i 12.8+2.2i
```

R Data frames

- a data frame is like a matrix, with a 2-dim rows-and-columns structure
- but each column may have a different mode
- as lists are the heterogeneous version of vectors, data frames are the heterogeneous analogs of matrices for two-dimensional data
- technically, a data frame is a list with all equal length vectors
- we can create a data frame combining two vectors

Exam ₁	Exam ₂	Gender
27	25	M
28	30	F
...		
27	27	M
25	28	F

```
> exam1 <- c(27,28,24,24,30,26,23,23,24,28,27,25)
> exam2 <- c(25,30,26,24,30,30,25,25,30,28,27,28)
> gender <- c("M","F","M","M","M","M","M","M","F","F","M","F")

> dc <- data.frame(exam1, exam2, gender)
> head(dc, n=2) # extract the first two lines of the data frame
  exam1 exam2 gender
1    27    25     M
2    28    30     F
```

Accessing data frames elements

- the data frame is a list, therefore we can access them via component index value `[[j]]` or via component names

```
> str(dc)
'data.frame':   12 obs. of  3 variables:
 $ exam1 : num  27 28 24 24 30 26 23 23 24 28 ...
 $ exam2 : num  25 30 26 24 30 30 25 25 30 28 ...
 $ gender: Factor w/ 2 levels "F","M": 2 1 2 2 2 2 2 2 1 1 ...

> dc[[1]] # access by component index
[1] 27 28 24 24 30 26 23 23 24 28 27 25

> dc$exam1 # access by component name
[1] 27 28 24 24 30 26 23 23 24 28 27 25
Levels: F M
```

- but a data frame can be treated in a matrix-like fashion, as well

```
> dc[,1] # select column 1
[1] 27 28 24 24 30 26 23 23 24 28 27 25

> dc[1,1] # and access the single element, as well
[1] 27
```


Advanced data frames : data selection

```
> dc[2:4,] # Select only rows 2:4
  exam1 exam2 gender
2     28     30      F
3     24     26      M
4     24     24      M
```

```
> dc[-(2:10),] # drop rows 2:10
  exam1 exam2 gender
1      27     25      M
11     27     27      M
12     25     28      F
```

- with the sample function , data can be selected at random

```
> dc[sample(1:12,3),] # select 3 rows at random
  exam1 exam2 gender
8      23     25      M
9      24     30      F
6      26     30      M
```

```
> dc[sample(1:12,3),] # select 3 rows at random
  exam1 exam2 gender
1      27     25      M
10     28     28      F
2      28     30      F
```

Advanced data frames : data selection

- suppose we want to extract all columns that contain numbers, rather than characters or logicals, from a data frame

```
> dc[,sapply(dc,is.numeric)]
  exam1 exam2
1      27     25
2      28     30
3      24     26
4      24     24
5      30     30
6      26     30
7      23     25
8      23     25
9      24     30
10     28     28
11     27     27
12     25     28
```

- and now we want to get only factors (and remove numerics)

```
> dc[,sapply(dc,is.factor)]
[1] M F M M M M M M F F M F
Levels: F M
```

Advanced data frames and NA elements

- sometimes our data frame can have missing values (NA) and we may need to omit those values
- we can create a shorter data frame using the `na.omit` function

```
> data
  slope pH area
1    11 4.1 3.6
2     NA 5.2 5.1
3     3 4.9 2.8
4     5  NA 3.7

> na.omit(data)
  slope pH area
1    11 4.1 3.6
3     3 4.9 2.8

> clean_data <- na.exclude(data)
> clean_data
  slope pH area
1    11 4.1 3.6
3     3 4.9 2.8

> lapply(clean_data, mean)
$ slope
[1] 7
$ pH
[1] 4.5
$ area
[1] 3.2

> # Let's count the missing values
> apply(apply(data, 2, is.na), 2, sum)
 slope pH area
    1   1   0
```

Advanced data frames : sorting elements

```
> dc[order(exam1),]
  exam1 exam2 gender
7     23    25     M
8     23    25     M
...
> dc[order(exam1, decreasing=TRUE),]
  exam1 exam2 gender
5     30    30     M
2     28    30     F

> dc[order(gender, exam1, exam2, decreasing=TRUE),]
  exam1 exam2 gender
5     30    30     M
11    27    27     M
1     27    25     M
6     26    30     M
3     24    26     M
4     24    24     M
7     23    25     M
8     23    25     M
2     28    30     F
10    28    28     F
12    25    28     F
9     24    30     F
```

Summary of data selection in data frames

- given a data frame called `data`, we assume `n` is a row number, and `m` is one of the column.
- the syntax `[n,]` selects all the columns given row `n`, while `[,m]` selects all the rows with column `m`

command	meaning
<code>data[n,]</code>	select all of the columns from row <code>n</code> of the data frame
<code>data[-n,]</code>	drop the whole of row <code>n</code> from the data frame
<code>data[1:n,]</code>	select all of the columns from rows 1 to <code>n</code> of the data frame
<code>data[-(1:n),]</code>	drop all of the columns from rows 1 to <code>n</code> of the data frame
<code>data[c(i,j,k),]</code>	select all of the columns from rows <code>i</code> , <code>j</code> , and <code>k</code> of the data frame
<code>data[x > y,]</code>	use a logical test ($x > y$) to select all columns from certain rows
<code>data[,m]</code>	select all of the rows from column <code>m</code> of the data frame
<code>data[, -m]</code>	drop the whole of column <code>m</code> from the data frame
<code>data[, 1:m]</code>	select all of the rows from columns 1 to <code>m</code> of the data frame
<code>data[, -(1:m)]</code>	drop all of the rows from columns 1 to <code>m</code> of the data frame
<code>data[, c(i,j,k)]</code>	select all of the rows from columns <code>i</code> , <code>j</code> , and <code>k</code> of the data frame
<code>data[, x > y]</code>	use a logical test ($x > y$) to select all rows from certain columns
<code>data[, c(1:m,i,j,k)]</code>	add duplicate copies of columns <code>i</code> , <code>j</code> , and <code>k</code> to the data frame
<code>data[x > y, a != b]</code>	extract certain rows ($x > y$) and certain columns ($a \neq b$)
<code>data[c(1:n,i,j,k),]</code>	add duplicate copies of rows <code>i</code> , <code>j</code> , and <code>k</code> to the data frame